

# ***AI BASED DIABETES PREDICTION SYSTEM***

**NAME : SABASTIN SELVIN JERALD J**

## **CONTEXT**

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective is to predict based on diagnostic measurements whether a patient has diabetes.

## **CONTENT**

Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

- Pregnancies: Number of times pregnant
- Glucose: Plasma glucose concentration a 2 hours in an oral glucose tolerance test
- Blood Pressure: Diastolic blood pressure (mm Hg)

- Skin Thickness: Triceps skin fold thickness (mm)
- Insulin: 2-Hour serum insulin (mu U/ml)
- BMI: Body mass index (weight in kg/(height in m)<sup>2</sup>)
- Diabetes Pedigree Function: Diabetes pedigree function
- Age: Age (years)
- Outcome: Class variable (0 or 1)

## PROBLEM STATEMENT

Develop an AI-powered diabetes prediction system that leverages machine learning algorithms to analyze medical data and predict the likelihood of an individual developing diabetes, providing early risk assessment and personalized preventive measures.

## PROBLEM DEFINITION

The problem is to build an AI-powered diabetes prediction system that uses machine learning algorithms to analyze medical data and predict the likelihood of an individual developing diabetes. The system aims to provide

early risk assessment and personalized preventive measures, allowing individuals to take proactive actions to manage their health.

## SOURCES

(a) Original owners: National Institute of Diabetes and Digestive and Kidney Diseases

(b) Donor of database: Vincent Sigillito  
(vgs@aplcn.apl.jhu.edu) Research Center, RMI Group  
Leader Applied Physics Laboratory The Johns Hopkins  
University Johns Hopkins Road Laurel, MD 20707 (301) 953-6231

(c) Date received: 9 May 1990

## DATA SET

<https://www.kaggle.com/datasets/mathchi/diabetes-data-set>

## DATA PRE PROCESSING

Gather a diverse and comprehensive dataset containing relevant features for diabetes prediction.- Preprocess the data, which may include tasks like normalization, handling missing values, and feature engineering.

## **ENSEMBLE LEARNING METHODS**

Use techniques like Random Forests, Gradient Boosting, or AdaBoost to create multiple models using subsets of the data.- Each model in the ensemble focuses on different aspects and learns from different parts of the data.

## **DEEP LEARNING ALGORITHM**

- Implement a deep learning model, perhaps a feedforward neural network or a recurrent neural network (RNN), to capture complex relationships within the data.

# CODING

```
[9]: import numpy as np
import pandas as pd
# import statsmodels.api as sm
import seaborn as sns
```

```
[7]: #Reading the dataset
df = pd.read_csv("diabetes.csv")
```

```
[8]: df.head()
```

```
[8]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0             6      148             72           35         0  33.6
1             1       85             66           29         0  26.6
2             8      183             64            0         0  23.3
3             1       89             66           23        94  28.1
4             0      137             40           35       168  43.1
```

```
DiabetesPedigreeFunction  Age  Outcome
0              0.627    50         1
1              0.351    31         0
2              0.672    32         1
3              0.167    21         0
4              2.288    33         1
```

```
[10]: # The size of the data set was examined. It consists of 768 observation units,
      ↪and 9 variables.
df.shape
```

```
[10]: (768, 9)
```

```
[11]: #Feature information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
#   Column              Non-Null Count  Dtype
```

```

0 Pregnancies          768 non-null    int64
1 Glucose              768 non-null    int64
2 BloodPressure        768 non-null    int64
3 SkinThickness        768 non-null    int64
4 Insulin              768 non-null    int64
5 BMI                  768 non-null    float64
6 DiabetesPedigreeFunction 768 non-null    float64
7 Age                  768 non-null    int64
8 Outcome              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```

[13]: # Descriptive statistics of the data set accessed.
df.describe([0.10,0.25,0.50,0.75,0.90,0.95,0.99]).T

```

```

[13]:

```

	count	mean	std	min	10%	\
Pregnancies	768.0	3.845052	3.369578	0.000	0.000	
Glucose	768.0	120.894531	31.972618	0.000	85.000	
BloodPressure	768.0	69.105469	19.355807	0.000	54.000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.000	
Insulin	768.0	79.799479	115.244002	0.000	0.000	
BMI	768.0	31.992578	7.884160	0.000	23.600	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.165	
Age	768.0	33.240885	11.760232	21.000	22.000	
Outcome	768.0	0.348958	0.476951	0.000	0.000	

	25%	50%	75%	90%	95%	\
Pregnancies	1.00000	3.0000	6.00000	9.0000	10.00000	
Glucose	99.00000	117.0000	140.25000	167.0000	181.00000	
BloodPressure	62.00000	72.0000	80.00000	88.0000	90.00000	
SkinThickness	0.00000	23.0000	32.00000	40.0000	44.00000	
Insulin	0.00000	30.5000	127.25000	210.0000	293.00000	
BMI	27.30000	32.0000	36.60000	41.5000	44.39500	
DiabetesPedigreeFunction	0.24375	0.3725	0.62625	0.8786	1.13285	
Age	24.00000	29.0000	41.00000	51.0000	58.00000	
Outcome	0.00000	0.0000	1.00000	1.0000	1.00000	

	99%	max
Pregnancies	13.00000	17.00
Glucose	196.00000	199.00
BloodPressure	106.00000	122.00
SkinThickness	51.33000	99.00
Insulin	519.90000	846.00
BMI	50.75900	67.10
DiabetesPedigreeFunction	1.69833	2.42
Age	67.00000	81.00
Outcome	1.00000	1.00


```
[14]: # The distribution of the Outcome variable was examined.  
df["Outcome"].value_counts()*100/len(df)
```

```
[14]: Outcome  
0    65.104167  
1    34.895833  
Name: count, dtype: float64
```

```
[15]: # The classes of the outcome variable were examined.  
df.Outcome.value_counts()
```

```
[15]: Outcome  
0    500  
1    268  
Name: count, dtype: int64
```

```
[16]: # The histogram of the Age variable was reached.  
df["Age"].hist(edgecolor = "black");
```



```
[18]: print("Max Age: " + str(df["Age"].max()) + " Min Age: " + str(df["Age"].min()))
```

```
Max Age: 81 Min Age: 21
```

```
[24]: # Histogram and density graphs of all variables were accessed.  
import matplotlib.pyplot as plt
```

```
# import statsmodels.api as sm  
fig, ax = plt.subplots(4,2, figsize=(16,16))  
sns.distplot(df.Age, bins = 20, ax=ax[0,0])  
sns.distplot(df.Pregnancies, bins = 20, ax=ax[0,1])  
sns.distplot(df.Glucose, bins = 20, ax=ax[1,0])  
sns.distplot(df.BloodPressure, bins = 20, ax=ax[1,1])  
sns.distplot(df.SkinThickness, bins = 20, ax=ax[2,0])  
sns.distplot(df.Insulin, bins = 20, ax=ax[2,1])  
sns.distplot(df.DiabetesPedigreeFunction, bins = 20, ax=ax[3,0])  
sns.distplot(df.BMI, bins = 20, ax=ax[3,1])
```

C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel\_9092\1084607558.py:6:  
UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.Age, bins = 20, ax=ax[0,0])
```

C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel\_9092\1084607558.py:7:  
UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.Pregnancies, bins = 20, ax=ax[0,1])
```

C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel\_9092\1084607558.py:8:  
UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see



<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.Glucose, bins = 20, ax=ax[1,0])
C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel_9092\1084607558.py:9:
UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.BloodPressure, bins = 20, ax=ax[1,1])
C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel_9092\1084607558.py:10:
UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.SkinThickness, bins = 20, ax=ax[2,0])
C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel_9092\1084607558.py:11:
UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.Insulin, bins = 20, ax=ax[2,1])
C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel_9092\1084607558.py:12:
UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.DiabetesPedigreeFunction, bins = 20, ax=ax[3,0])
```

C:\Users\AJITHKUMAR\AppData\Local\Temp\ipykernel\_9092\1084607558.py:13:  
UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.BMI, bins = 20, ax=ax[3,1])
```

[24]: <Axes: xlabel='BMI', ylabel='Density'>



```
[25]: df.groupby("Outcome").agg({"Pregnancies": "mean"})
```

```
[25]:      Pregnancies
Outcome
0          3.298000
1          4.865672
```

```
[26]: df.groupby("Outcome").agg({"Age": "mean"})
```

```
[26]:      Age
Outcome
0      31.190000
1      37.067164
```

```
[27]: df.groupby("Outcome").agg({"Age": "max"})
```

```
[27]:      Age
Outcome
0      81
1      70
```

```
[28]: df.groupby("Outcome").agg({"Insulin": "mean"})
```

```
[28]:      Insulin
Outcome
0      68.792000
1     100.335821
```

```
[29]: df.groupby("Outcome").agg({"Insulin": "max"})
```

```
[29]:      Insulin
Outcome
0      744
1      846
```

```
[30]: df.groupby("Outcome").agg({"Glucose": "mean"})
```

```
[30]:      Glucose
Outcome
0     109.980000
1     141.257463
```

```
[31]: df.groupby("Outcome").agg({"Glucose": "max"})
```

```
[31]:      Glucose
Outcome
0      197
1      199
```

```
[32]: df.groupby("Outcome").agg({"BMI": "mean"})
```


```
[32]:      BMI
Outcome
0     30.304200
1     35.142537
```

```
[33]: # The distribution of the outcome variable in the data was examined and
      ↪ visualized.
      f,ax=plt.subplots(1,2,figsize=(18,8))
      df['Outcome'].value_counts().plot.pie(explode=[0,0.1],autopct='%1.
      ↪1f%%',ax=ax[0],shadow=True)
```

```
ax[0].set_title('target')
ax[0].set_ylabel('')
sns.countplot('Outcome',data=df,ax=ax[1])
ax[1].set_title('Outcome')
plt.show()
```

```
-----
TypeError                                Traceback (most recent call last)
Cell In[33], line 6
      4 ax[0].set_title('target')
      5 ax[0].set_ylabel('')
----> 6 sns.countplot('Outcome',data=df,ax=ax[1])
      7 ax[1].set_title('Outcome')
      8 plt.show()
```

TypeError: countplot() got multiple values for argument 'data'



[34]: df.corr()

```
[34]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	

Outcome	0.221898	0.466581	0.065068	0.074752
---------	----------	----------	----------	----------

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683		-0.033523
Glucose	0.331357	0.221071		0.137337
BloodPressure	0.088933	0.281805		0.041265
SkinThickness	0.436783	0.392573		0.183928
Insulin	1.000000	0.197859		0.185071
BMI	0.197859	1.000000		0.140647
DiabetesPedigreeFunction	0.185071	0.140647		1.000000
Age	-0.042163	0.036242		0.033561
Outcome	0.130548	0.292695		0.173844

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[35]: # Correlation matrix graph of the data set
f, ax = plt.subplots(figsize= [20,15])
sns.heatmap(df.corr(), annot=True, fmt=".2f", ax=ax, cmap = "magma" )
ax.set_title("Correlation Matrix", fontsize=20)
plt.show()
```



```
[36]: df[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] =
      ~df[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']].replace(0, np.
      ~NaN)
      df.head()
```

```
[36]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI \
0           6    148.0         72.0         35.0      NaN  33.6
1           1     85.0         66.0         29.0      NaN  26.6
2           8    183.0         64.0          NaN      NaN  23.3
3           1     89.0         66.0         23.0    94.0  28.1
4           0    137.0         40.0         35.0   168.0  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
3                0.167    21         0
```



4 2.288 33 1


```
[37]: # Now, we can look at where are missing values
df.isnull().sum()
```

```
[37]: Pregnancies      0
      Glucose          5
      BloodPressure    35
      SkinThickness    227
      Insulin          374
      BMI              11
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[45]: from sklearn.preprocessing import scale, StandardScaler
      from sklearn.model_selection import train_test_split, GridSearchCV,
      ↪cross_val_score
      from sklearn.metrics import confusion_matrix, accuracy_score,
      ↪mean_squared_error, r2_score, roc_auc_score, roc_curve, classification_report
      from sklearn.linear_model import LogisticRegression
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.svm import SVC
      from sklearn.neural_network import MLPClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.ensemble import GradientBoostingClassifier
      from lightgbm import LGBMClassifier
      from sklearn.model_selection import KFold
      import warnings
      warnings.simplefilter(action = "ignore")
```

```
[47]: # Have been visualized using the missingno library for the visualization of
      ↪missing observations.
      # Plotting
      import missingno as msno
      msno.bar(df);
```





```
[48]: # The missing values will be filled with the median values of each variable.
def median_target(var):
    temp = df[df[var].notnull()]
    temp = temp[[var, 'Outcome']].groupby(['Outcome'])[var].median().
    reset_index()
    return temp

[49]: # The values to be given for incomplete observations are given the median value
      of people who are not sick and the median values of people who are sick.
columns = df.columns
columns = columns.drop("Outcome")
for i in columns:
    median_target(i)
    df.loc[(df['Outcome'] == 0) & (df[i].isnull()), i] = median_target(i)[0]
    df.loc[(df['Outcome'] == 1) & (df[i].isnull()), i] = median_target(i)[1]

[50]: df.head()
```

```
[50]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI \
0           6    148.0           72.0           35.0    169.5  33.6
1           1     85.0           66.0           29.0    102.5  26.6
2           8    183.0           64.0           32.0    169.5  23.3
3           1     89.0           66.0           23.0     94.0  28.1
4           0    137.0           40.0           35.0    168.0  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
```

1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[51]: # Missing values were filled.
df.isnull().sum()
```

```
[51]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[52]: # In the data set, there were asked whether there were any outlier observations.
      ↳ compared to the 25% and 75% quarters.
      # It was found to be an outlier observation.
      for feature in df:


          Q1 = df[feature].quantile(0.25)
          Q3 = df[feature].quantile(0.75)
          IQR = Q3-Q1
          lower = Q1- 1.5*IQR
          upper = Q3 + 1.5*IQR

          if df[(df[feature] > upper)].any(axis=None):
              print(feature,"yes")
          else:
              print(feature, "no")
```

```
Pregnancies yes
Glucose no
BloodPressure yes
SkinThickness yes
Insulin yes
BMI yes
DiabetesPedigreeFunction yes
Age yes
Outcome no
```


```
[53]: # The process of visualizing the Insulin variable with boxplot method was done.
      ↳ We find the outlier observations on the chart.
```

```
import seaborn as sns
sns.boxplot(x = df["Insulin"]);
```



```
[54]: #We conduct a stand alone observation review for the Insulin variable
      #We suppress contradictory values
      Q1 = df.Insulin.quantile(0.25)
      Q3 = df.Insulin.quantile(0.75)
      IQR = Q3-Q1
      lower = Q1 - 1.5*IQR
      upper = Q3 + 1.5*IQR
      df.loc[df["Insulin"] > upper, "Insulin"] = upper
```

```
[55]: import seaborn as sns
      sns.boxplot(x = df["Insulin"]);
```



```
[56]: # We determine outliers between all variables with the LOF method
from sklearn.neighbors import LocalOutlierFactor
lof = LocalOutlierFactor(n_neighbors= 10)
lof.fit_predict(df)
```

```
[56]: array([ 1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1, -1,  1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  
              1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1])
```

```
[57]: df_scores = lof.negative_outlier_factor_
      np.sort(df_scores)[0:30]
```

```
[58]: #We choose the threshold value according to lof scores
threshold = np.sort(df_scores)[7]
threshold
```

```
[59]: #We delete those that are higher than the threshold
outlier = df_scores > threshold
df = df[outlier]
# The size of the data set was examined.
df.shape
```

```
[59]: (760, 9)
```

```
[60]: # According to BMI, some ranges were determined and categorical variables were
      ↪ assigned.
NewBMI = pd.Series(["Underweight", "Normal", "Overweight", "Obesity 1",
      ↪ "Obesity 2", "Obesity 3"], dtype = "category")
df["NewBMI"] = NewBMI
df.loc[df["BMI"] < 18.5, "NewBMI"] = NewBMI[0]
df.loc[(df["BMI"] > 18.5) & (df["BMI"] <= 24.9), "NewBMI"] = NewBMI[1]
df.loc[(df["BMI"] > 24.9) & (df["BMI"] <= 29.9), "NewBMI"] = NewBMI[2]
df.loc[(df["BMI"] > 29.9) & (df["BMI"] <= 34.9), "NewBMI"] = NewBMI[3]
df.loc[(df["BMI"] > 34.9) & (df["BMI"] <= 39.9), "NewBMI"] = NewBMI[4]
df.loc[df["BMI"] > 39.9, "NewBMI"] = NewBMI[5]
```

```
[61]: df.head()
```

```
[61]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI \
0             6    148.0             72.0             35.0    169.5  33.6
1             1     85.0             66.0             29.0    102.5  26.6
2             8    183.0             64.0             32.0    169.5  23.3
3             1     89.0             66.0             23.0     94.0  28.1
4             0    137.0             40.0             35.0    168.0  43.1

      DiabetesPedigreeFunction  Age  Outcome  NewBMI
0                0.627    50         1  Obesity 1
1                0.351    31         0  Overweight
2                0.672    32         1   Normal
3                0.167    21         0  Overweight
4                2.288    33         1  Obesity 3
```

```
[62]: # A categorical variable creation process is performed according to the insulin
      ↪ value.
def set_insulin(row):
    if row["Insulin"] >= 16 and row["Insulin"] <= 166:
        return "Normal"
    else:
        return "Abnormal"
```

```
[63]: # The operation performed was added to the dataframe.
df = df.assign(NewInsulinScore=df.apply(set_insulin, axis=1))
```



```
df.head()
```

```
[63]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6      148.0           72.0           35.0    169.5   33.6
1           1       85.0           66.0           29.0    102.5   26.6
2           8      183.0           64.0           32.0    169.5   23.3
3           1       89.0           66.0           23.0     94.0   28.1
4           0      137.0           40.0           35.0    168.0   43.1

      DiabetesPedigreeFunction  Age  Outcome  NewBMI  NewInsulinScore
0                0.627    50         1  Obesity 1          Abnormal
1                0.351    31         0  Overweight          Normal
2                0.672    32         1    Normal          Abnormal
3                0.167    21         0  Overweight          Normal
4                2.288    33         1  Obesity 3          Abnormal
```

```
[64]: # Some intervals were determined according to the glucose variable and these
      were assigned categorical variables.
NewGlucose = pd.Series(["Low", "Normal", "Overweight", "Secret", "High"], dtype=
      "category")
df["NewGlucose"] = NewGlucose
df.loc[df["Glucose"] <= 70, "NewGlucose"] = NewGlucose[0]
df.loc[(df["Glucose"] > 70) & (df["Glucose"] <= 99), "NewGlucose"] =
      NewGlucose[1]
df.loc[(df["Glucose"] > 99) & (df["Glucose"] <= 126), "NewGlucose"] =
      NewGlucose[2]
df.loc[df["Glucose"] > 126, "NewGlucose"] = NewGlucose[3]
```

```
[65]: df.head()
```

```
[65]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6      148.0           72.0           35.0    169.5   33.6
1           1       85.0           66.0           29.0    102.5   26.6
2           8      183.0           64.0           32.0    169.5   23.3
3           1       89.0           66.0           23.0     94.0   28.1
4           0      137.0           40.0           35.0    168.0   43.1

      DiabetesPedigreeFunction  Age  Outcome  NewBMI  NewInsulinScore  \
0                0.627    50         1  Obesity 1          Abnormal
1                0.351    31         0  Overweight          Normal
2                0.672    32         1    Normal          Abnormal
3                0.167    21         0  Overweight          Normal
4                2.288    33         1  Obesity 3          Abnormal

      NewGlucose
0      Secret
1      Normal
```

```

2     Secret
3     Normal
4     Secret

```

```

[66]: # Here, by making One Hot Encoding transformation, categorical variables were
      ↪ converted into numerical values. It is also protected from the Dummy
      ↪ variable trap.
      df = pd.get_dummies(df, columns=["NewBMI", "NewInsulinScore", "NewGlucose"],
      ↪ drop_first = True)

```

```

[67]: df.head()

```

```

[67]:  Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0         6    148.0         72.0         35.0    169.5   33.6
1         1     85.0         66.0         29.0    102.5   26.6
2         8    183.0         64.0         32.0    169.5   23.3
3         1     89.0         66.0         23.0     94.0   28.1
4         0    137.0         40.0         35.0    168.0   43.1

      DiabetesPedigreeFunction  Age  Outcome  NewBMI_Obesity 1  NewBMI_Obesity 2  \
0                0.627    50         1         True         False
1                0.351    31         0         False         False
2                0.672    32         1         False         False
3                0.167    21         0         False         False
4                2.288    33         1         False         False

      NewBMI_Obesity 3  NewBMI_Overweight  NewBMI_Underweight  \
0                False                False                False
1                False                 True                False
2                False                False                False
3                False                 True                False
4                 True                False                False

      NewInsulinScore_Normal  NewGlucose_Low  NewGlucose_Normal  \
0                False                False                False
1                 True                False                 True
2                False                False                False
3                 True                False                 True
4                False                False                False

      NewGlucose_Overweight  NewGlucose_Secret
0                False                 True
1                False                False
2                False                 True
3                False                False
4                False                 True

```



```
[68]: categorical_df = df[['NewBMI_Obesity 1', 'NewBMI_Obesity 2', 'NewBMI_Obesity 3',
    ↪ 'NewBMI_Overweight', 'NewBMI_Underweight',
    ↪
    ↪ 'NewInsulinScore_Normal', 'NewGlucose_Low', 'NewGlucose_Normal',
    ↪ 'NewGlucose_Overweight', 'NewGlucose_Secret']]
categorical_df.head()
```

```
[68]:   NewBMI_Obesity 1  NewBMI_Obesity 2  NewBMI_Obesity 3  NewBMI_Overweight \
0                True                False                False                False
1                False                False                False                True
2                False                False                False                False
3                False                False                False                True
4                False                False                True                False

   NewBMI_Underweight  NewInsulinScore_Normal  NewGlucose_Low \
0                False                False                False
1                False                True                False
2                False                False                False
3                False                True                False
4                False                False                False

   NewGlucose_Normal  NewGlucose_Overweight  NewGlucose_Secret
0                False                False                True
1                True                False                False
2                False                False                True
3                True                False                False
4                False                False                True
```

```
[69]: y = df["Outcome"]
X = df.drop(["Outcome", 'NewBMI_Obesity 1', 'NewBMI_Obesity 2', 'NewBMI_Obesity_
    ↪ 3', 'NewBMI_Overweight', 'NewBMI_Underweight',
    ↪
    ↪ 'NewInsulinScore_Normal', 'NewGlucose_Low', 'NewGlucose_Normal',
    ↪ 'NewGlucose_Overweight', 'NewGlucose_Secret'], axis = 1)
cols = X.columns
index = X.index
```

```
[70]: X.head()
```

```
[70]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI \
0           6    148.0           72.0           35.0    169.5  33.6
1           1     85.0           66.0           29.0    102.5  26.6
2           8    183.0           64.0           32.0    169.5  23.3
3           1     89.0           66.0           23.0     94.0  28.1
4           0    137.0           40.0           35.0    168.0  43.1

DiabetesPedigreeFunction  Age
```

0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33

```
[71]: # The variables in the data set are an effective factor in increasing the
      ↪ performance of the models by standardization.
      # There are multiple standardization methods. These are methods such as "
      ↪ Normalize", " MinMax", " Robust" and "Scale".
      from sklearn.preprocessing import RobustScaler
      transformer = RobustScaler().fit(X)
      X = transformer.transform(X)
      X = pd.DataFrame(X, columns = cols, index = index)
```

```
[72]: X.head()
```

```
[72]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin    BMI  \
0          0.6    0.775         0.000      1.000000  1.000000  0.177778
1         -0.4   -0.800        -0.375      0.142857  0.000000 -0.600000
2          1.0    1.650        -0.500      0.571429  1.000000 -0.966667
3         -0.4   -0.700        -0.375     -0.714286 -0.126866 -0.433333
4         -0.6    0.500        -2.000      1.000000  0.977612  1.233333

      DiabetesPedigreeFunction  Age
0          0.669707  1.235294
1         -0.049511  0.117647
2          0.786971  0.176471
3         -0.528990 -0.470588
4          4.998046  0.235294
```

```
[73]: X = pd.concat([X,categorical_df], axis = 1)
```

```
[74]: X.head()
```

```
[74]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin    BMI  \
0          0.6    0.775         0.000      1.000000  1.000000  0.177778
1         -0.4   -0.800        -0.375      0.142857  0.000000 -0.600000
2          1.0    1.650        -0.500      0.571429  1.000000 -0.966667
3         -0.4   -0.700        -0.375     -0.714286 -0.126866 -0.433333
4         -0.6    0.500        -2.000      1.000000  0.977612  1.233333

      DiabetesPedigreeFunction  Age  NewBMI_Obesity 1  NewBMI_Obesity 2  \
0          0.669707  1.235294             True             False
1         -0.049511  0.117647             False             False
2          0.786971  0.176471             False             False
3         -0.528990 -0.470588             False             False
```

4	4.998046	0.235294	False	False
---	----------	----------	-------	-------

	NewBMI_Obesity 3	NewBMI_Overweight	NewBMI_Underweight	\
0	False	False	False	
1	False	True	False	
2	False	False	False	
3	False	True	False	
4	True	False	False	

	NewInsulinScore_Normal	NewGlucose_Low	NewGlucose_Normal	\
0	False	False	False	
1	True	False	True	
2	False	False	False	
3	True	False	True	
4	False	False	False	

	NewGlucose_Overweight	NewGlucose_Secret
0	False	True
1	False	False
2	False	True
3	False	False
4	False	True

```
[75]: y.head()
```

```
[75]: 0    1
      1    0
      2    1
      3    0
      4    1
      Name: Outcome, dtype: int64
```

```
[ ]: # Validation scores of all base models

models = []
models.append(('LR', LogisticRegression(random_state = 12345)))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier(random_state = 12345)))
models.append(('RF', RandomForestClassifier(random_state = 12345)))
models.append(('SVM', SVC(gamma='auto', random_state = 12345)))
models.append(('XGB', GradientBoostingClassifier(random_state = 12345)))
models.append(('LightGBM', LGBMClassifier(random_state = 12345)))

# evaluate each model in turn
results = []
names = []
```

```
[ ]: for name, model in models:

    kfold = KFold(n_splits = 10, random_state = 12345)
    cv_results = cross_val_score(model, X, y, cv = 10, scoring= "accuracy")
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

# boxplot algorithm comparison
fig = plt.figure(figsize=(15,10))
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```

## BENEFITS

**1. Early Detection:** AI models can analyze large datasets of patient information, including medical records, genetics, lifestyle data, and more, to detect patterns and risk factors associated with diabetes. This early detection can help identify individuals at risk before symptoms become severe.

**2. Personalized Risk Assessment:** AI can provide personalized risk assessments by taking into account an individual's unique health history and factors, such as age, family history, and lifestyle. This allows for tailored interventions and preventive measures.

**3. Predictive Accuracy:** AI models have the ability to process and analyze vast amounts of data, leading to more accurate predictions. These predictions can help healthcare professionals make informed decisions and recommendations for patients.

**4. Cost-Effective:** Early detection and prevention of diabetes can lead to cost savings in healthcare by reducing the need for expensive treatments and hospitalizations. This can lower the overall economic burden of diabetes on healthcare systems.

**5. Continuous Monitoring:** AI models can be integrated with wearable devices and mobile apps to provide continuous monitoring of patients' health. This real-time data can alert individuals and healthcare providers to changes in their risk factors and overall health.

**6. Improved Healthcare Delivery:** Healthcare providers can use AI-based prediction models to prioritize patients at higher risk of developing diabetes, ensuring they receive necessary interventions and education. This can lead to more efficient and effective healthcare delivery.

**7. Research and Insights:** AI can help researchers analyze large-scale population data to identify new risk factors,

contributing to a better understanding of diabetes and potential new avenues for prevention and treatment.

**8. Patient Empowerment:** By providing individuals with information about their diabetes risk, AI-based models empower them to make informed decisions about their health and take proactive steps to prevent or manage the condition.

**9. Reduction of Misdiagnosis:** AI can assist in reducing misdiagnosis or delayed diagnosis by considering a wide range of patient data and reducing the impact of human bias in the diagnostic process.

**10. Scalability:** AI models can be deployed at scale, making them accessible to a larger population and ensuring that more people can benefit from early diabetes prediction and prevention.

## CONCLUSION

However, as we embrace the potential of AI in diabetes prediction, it is crucial to consider ethical concerns, data privacy, and the need for a human-centered approach. AI should complement, rather than replace, the expertise of healthcare professionals. Furthermore, ensuring that these models are accessible and equitable is of paramount importance to maximize their positive impact.

In the years to come, the continued development and integration of AI-based diabetes prediction models will likely play a crucial role in

reducing the burden of diabetes, improving patient outcomes, and advancing our understanding of this complex condition. As technology evolves and research progresses, we can anticipate even more refined and effective tools to tackle diabetes, ultimately leading to better health and well-being for individuals worldwide.

PRESENTED BY,

**SABASTIN SELVIN JERALD J**